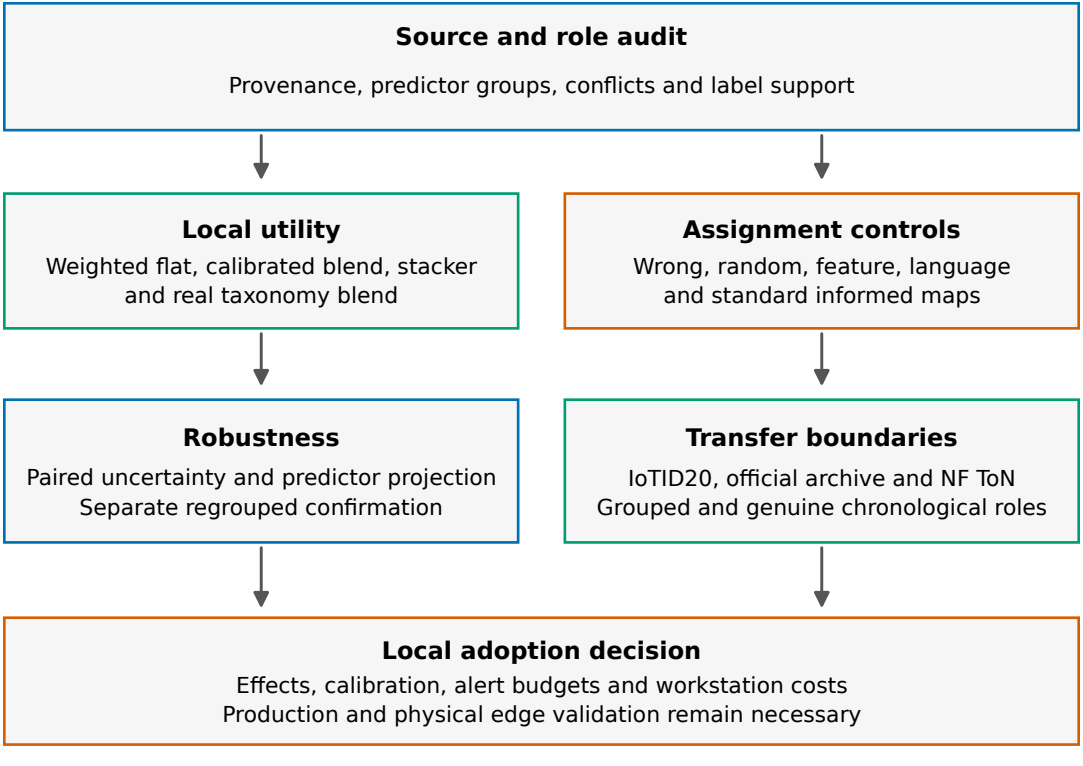
Source and role audit
Provenance, predictor groups, conflicts and label support
Local utility
Weighted flat, calibrated blend, stacker and real taxonomy blend
Assignment controls
Wrong, random, feature, language and standard informed maps
Robustness
Paired uncertainty and predictor projection
Separate regrouped confirmation
Transfer boundaries
IoTID20, official archive and NF ToN
Grouped and genuine chronological roles
Local adoption decision
Effects, calibration, alert budgets and workstation costs
Production and physical edge validation remain necessary